SCHEMA 4:

1)A)OLD QUERY 10 WITH NO INDEXES:

Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =off;
3-set enable_indexscan = off;
4- set enable_indexonlyscan = off;

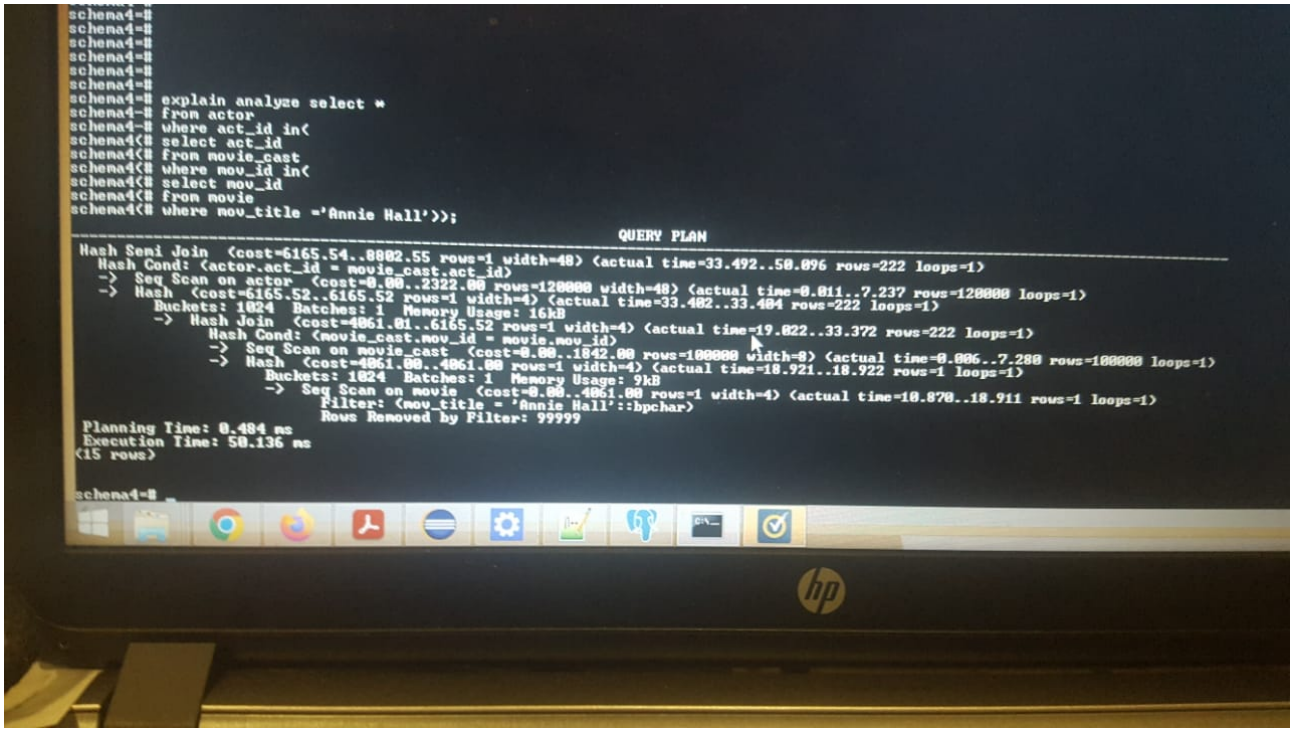

**********CONCLUSION OLD QUERY 10 NO INDEXES:**********

1-OLD QUERY 10 NO INDEXES COST = 8802
2-OLD QUERY 10 NO INDEXES TIME = 50 MS

1)B)OLD QUERY 10 WITH B+ INDEXES:

B+Indexes:
1- All default PKs indexes.
2 - movie(mov_title)
3 - movie_cast(mov_id).
Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

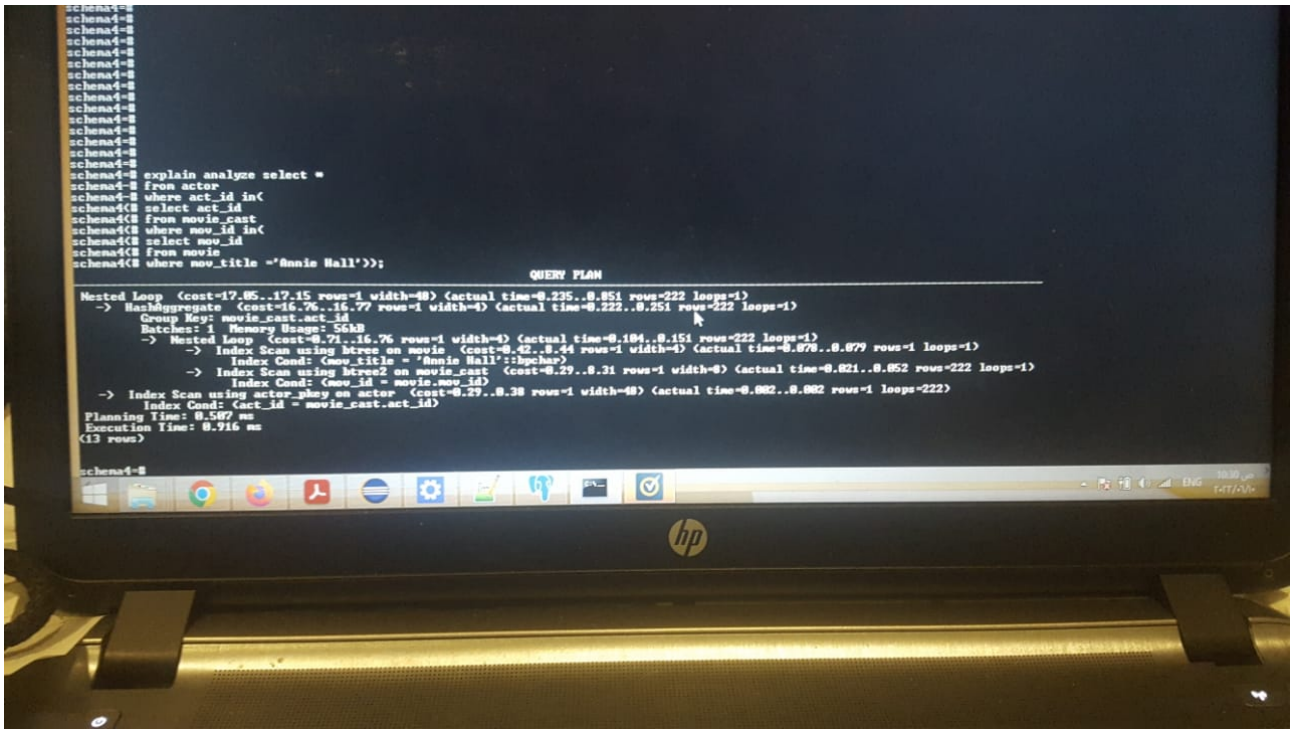

**********CONCLUSION QUERY 10 B+ INDEXES:**********
1-OLD QUERY 10 NO INDEXES COST = 8802
2-OLD QUERY 10 NO INDEXES TIME = 50 MS
————————————————————————————
1-OLD QUERY 10 B+ INDEXES COST = 17
2-OLD QUERY 10 B+ INDEXES TIME = 0.9 MS
=> 99.8% more cost efficient & 98.2% more time efficient.

1)C)OLD QUERY 10 WITH Hash INDEXES:

Hash indexes:1-actor(act_id)
2-movie_cast(mov_id)
3-movie(mov_title)

Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

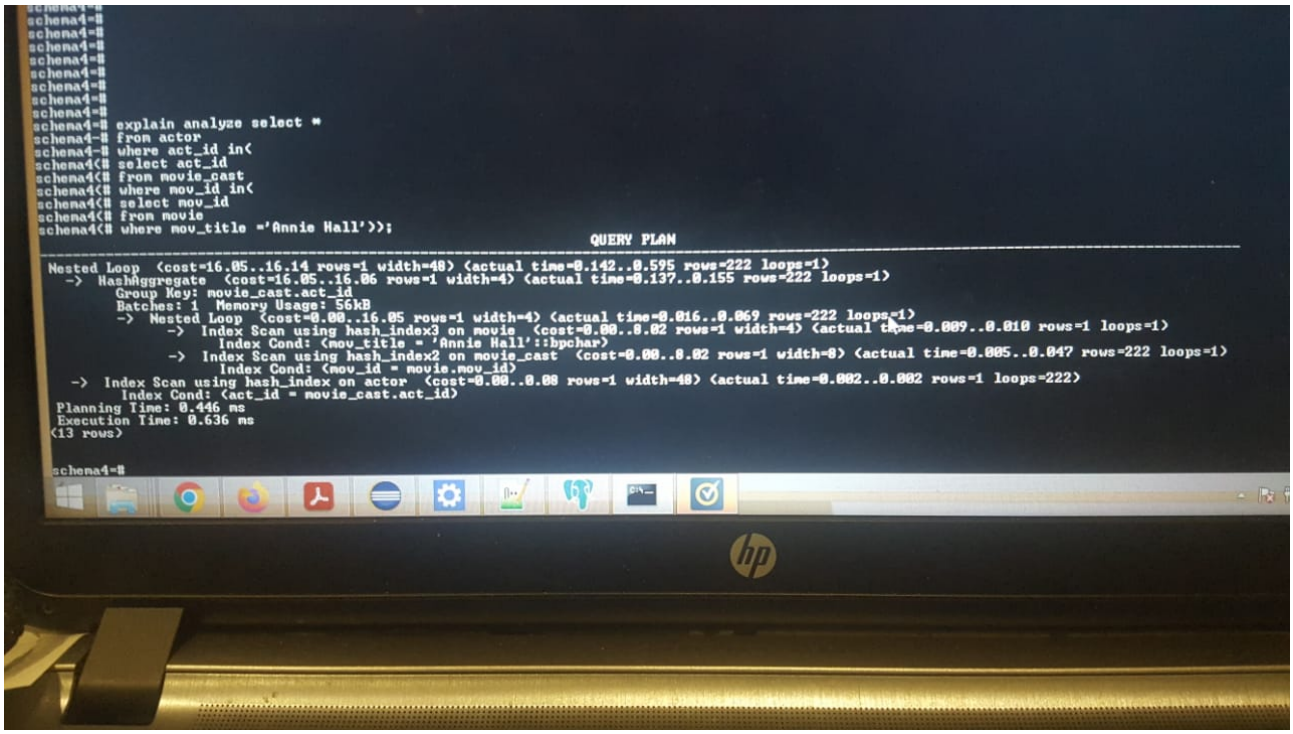

**********CONCLUSION QUERY 10 HASH INDEXES:**********

1-OLD QUERY 10 NO INDEXES COST =8802
2-OLD QUERY 10 NO INDEXES TIME = 50 MS

————————————————————————

1-OLD QUERY 10 HASH INDEXES COST = 16
2-OLD QUERY 10 HASH INDEXES TIME = 0.6 MS
=>99.8% more cost efficient &98.8% more time efficient.

1)D)OLD QUERY 10 WITH BRIN INDEX:

BRIN INDEXES: 1-movie(mov_title)
2- movie_cast(mov_id)

Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;
5- update pg_index set indisvalid = false where indexrelid = 'actor_pkey'::regclass ;

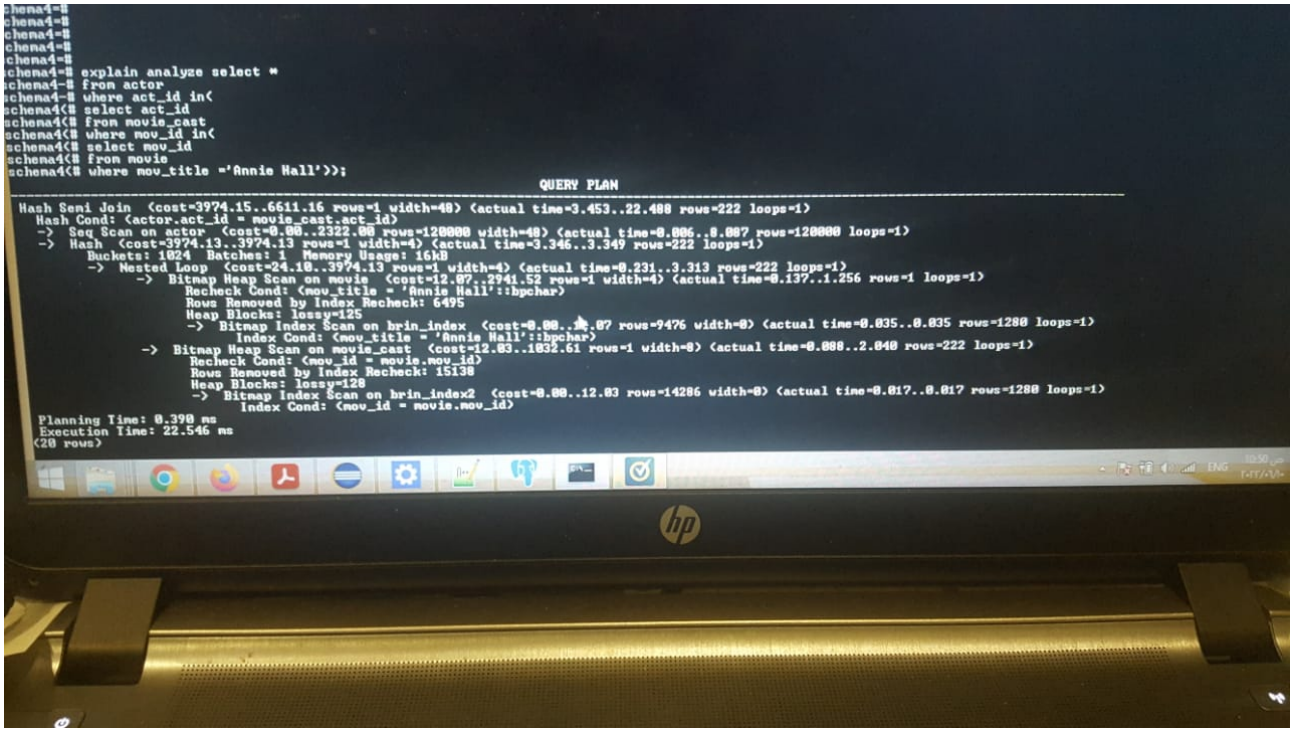

**********CONCLUSION QUERY 10 BRIN INDEX:**********

1-OLD QUERY 10 NO INDEXES COST = 8802
2-OLD QUERY 10 NO INDEXES TIME = 50 MS

1-OLD QUERY 10 BRIN INDEX COST = 6611
2-OLD QUERY 10 BRIN INDEX TIME = 22 MS
=>25% more cost efficient &56% more time efficient.

1)E)OLD QUERY 10 WITH MY CHOICE FOR INDEXES:

Full hash indexes same as 1)c): as all the conditions are equality conditions
Hash indexes:1-actor(act_id)
2-movie_cast(mov_id)
3-movie(mov_title)
Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

**********CONCLUSION QUERY 10 HASH INDEXES:**********
1-OLD QUERY 10 NO INDEXES COST =8802
2-OLD QUERY 10 NO INDEXES TIME = 50 MS
————————————————————————————————
1-OLD QUERY 10 HASH INDEXES COST = 16
2-OLD QUERY 10 HASH INDEXES TIME = 0.6 MS
=>99.8% more cost efficient &98.8% more time efficient.

2)A)OPTIMIZED QUERY 10 WITH NO INDEXES:

Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =off;
3-set enable_indexscan = off;
4- set enable_indexonlyscan = off;

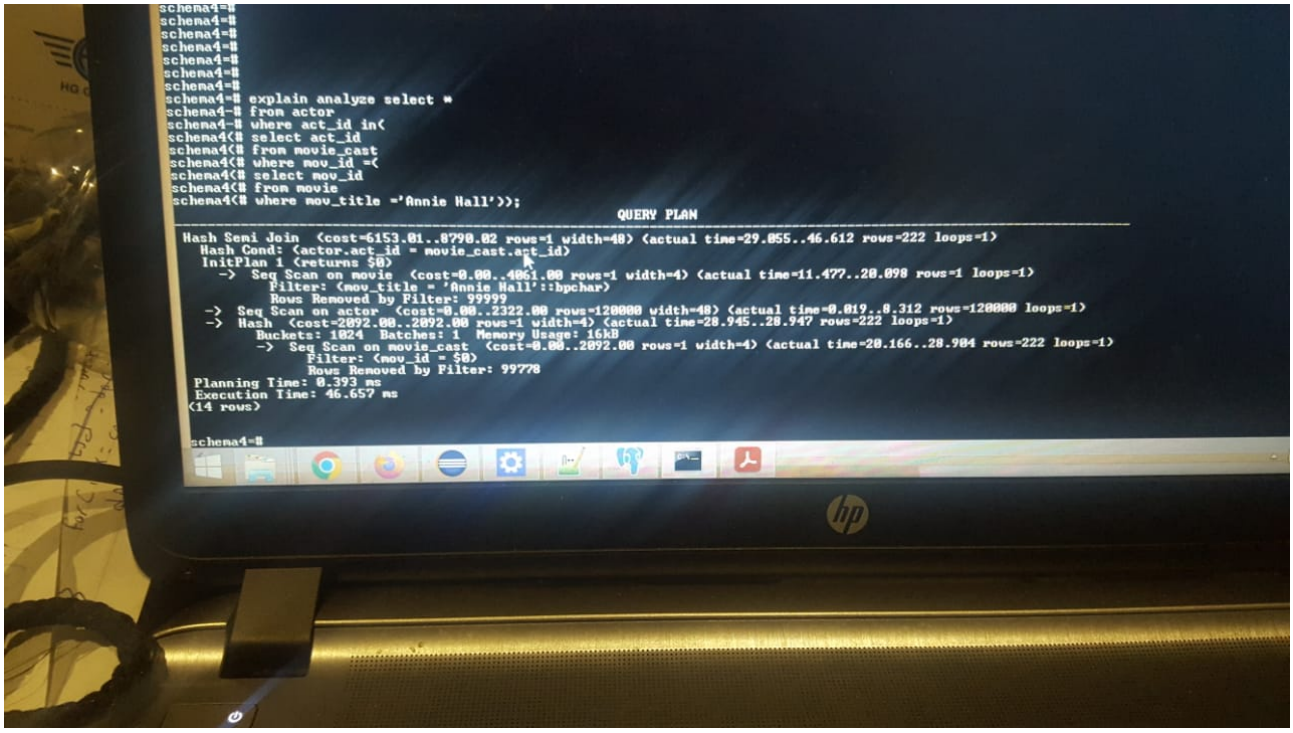

Optimised query 10 no index cost = 6153
Optimised query 10 no index time = 46.6

2)B)OPTIMIZED QUERY 10 WITH B+ INDEXES:

B+Indexes:
1- All default PKs indexes.
2 - movie(mov_title)
3 - movie_cast(mov_id).
Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

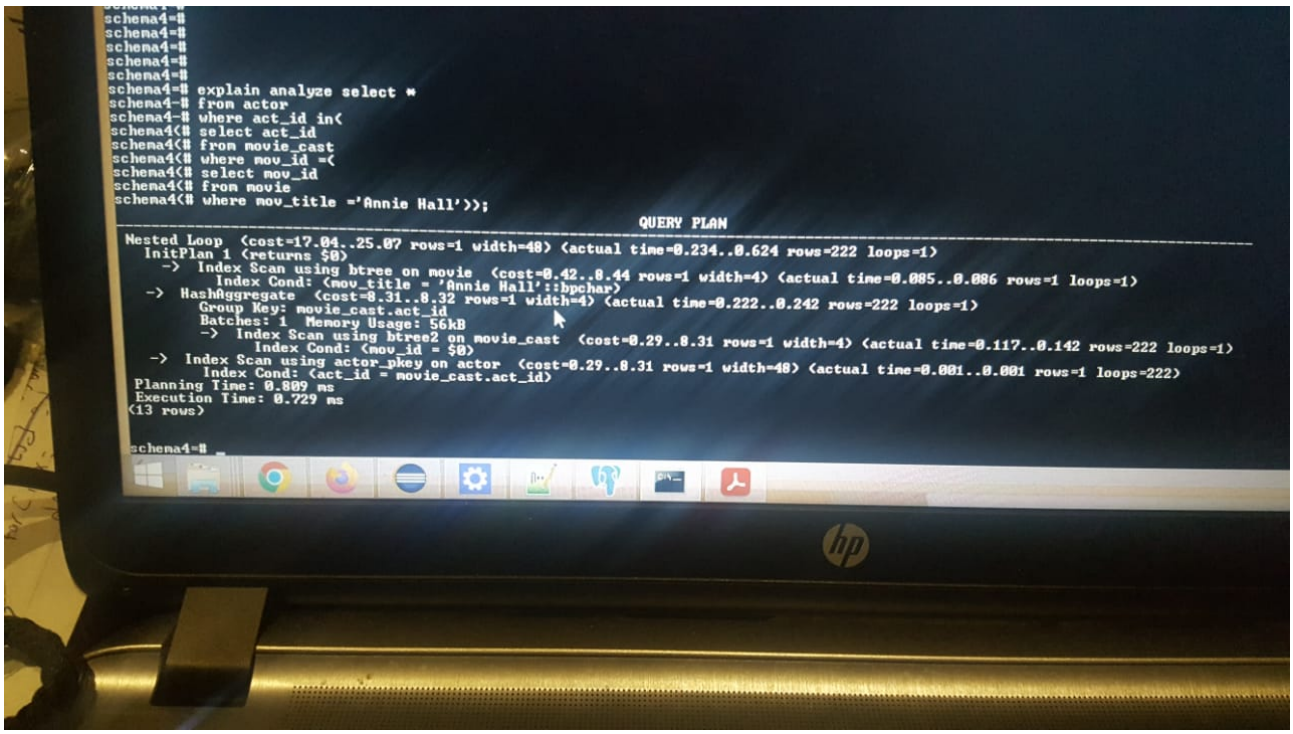

Cost = 25
Time =0.7ms

1)C)OPTIMIZED QUERY 10 WITH Hash INDEXES:

Hash indexes:1-actor(act_id)
2-movie_cast(mov_id)
3-movie(mov_title)

Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

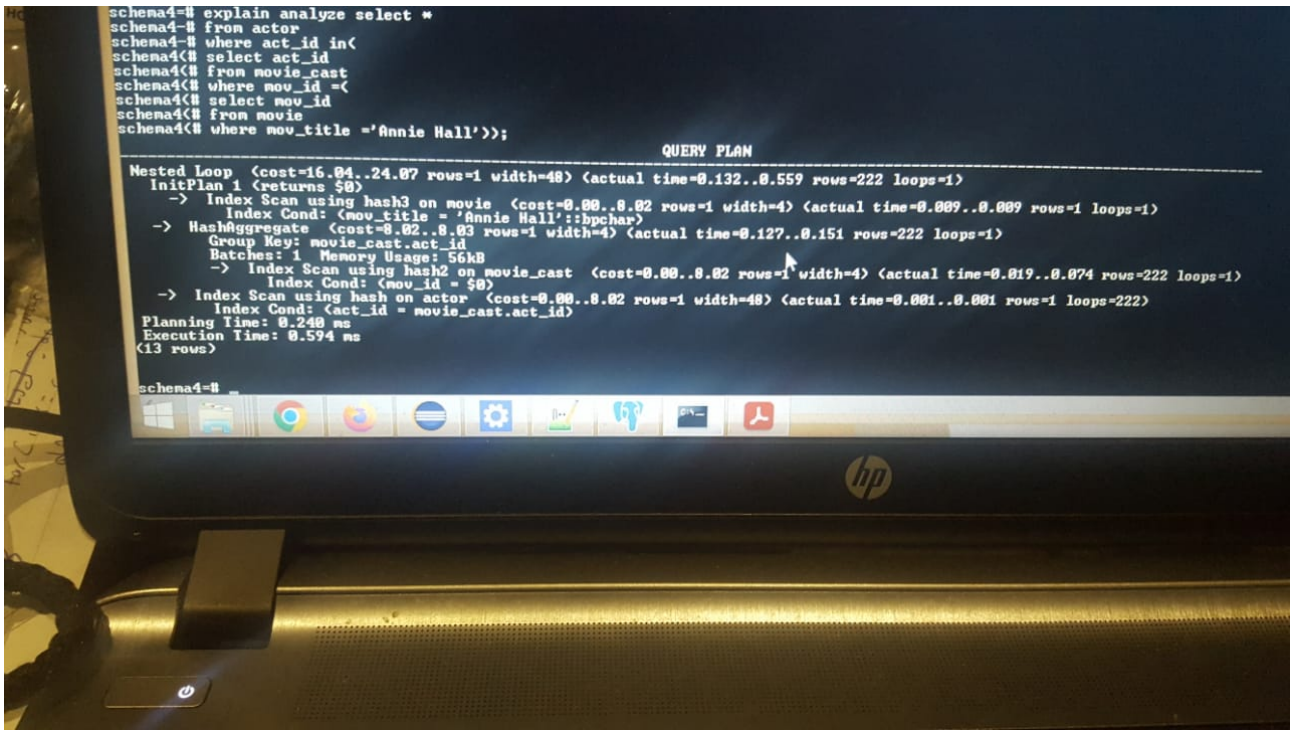

New cost hash = 24
New time hash = .2MS

1)D)optimized QUERY 10 WITH BRIN INDEX:

BRIN INDEXES: 1-movie(mov_title)
2- movie_cast(mov_id)
Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;
5- update pg_index set indisvalid = false where indexrelid = 'actor_pkey'::regclass ;

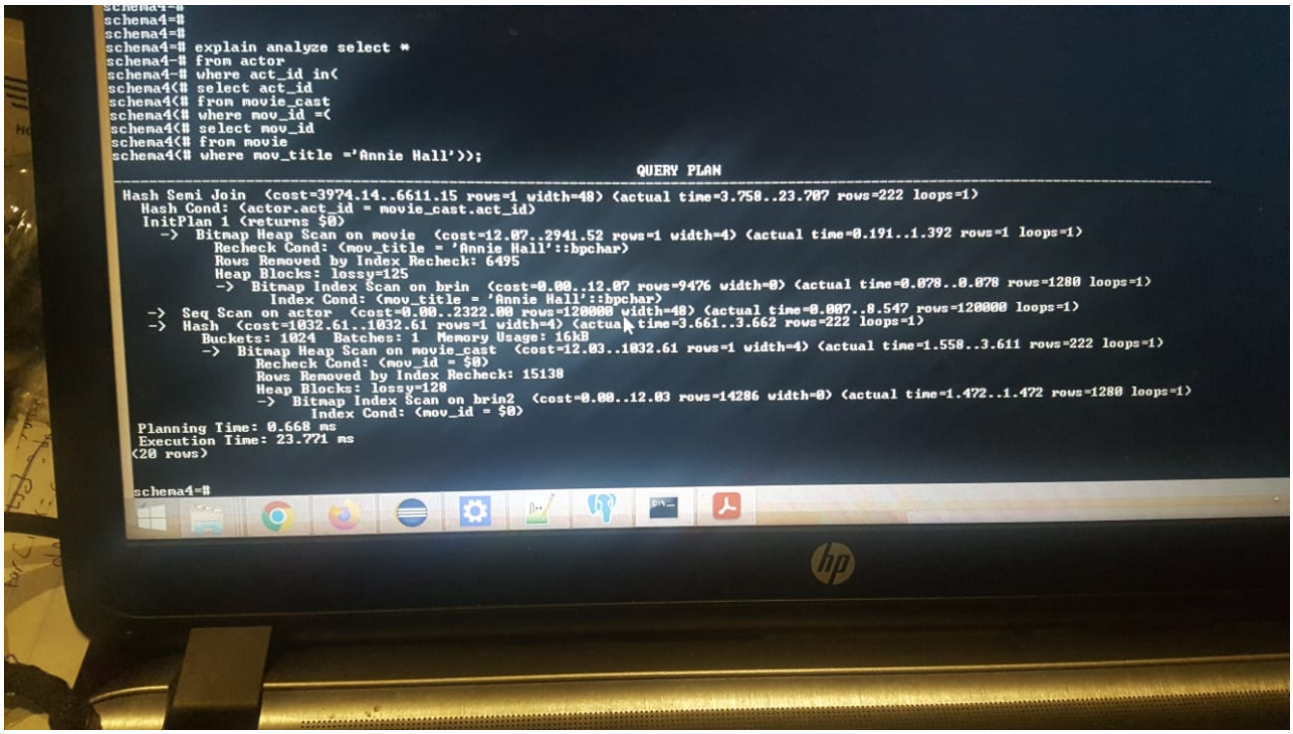

Cost = 6611
Time =23.7 ms

1)E)OPTIMIZED QUERY 10 WITH MY CHOICE FOR INDEXES:

Full hash indexes same as 1)c): as all the conditions are equality conditions
Hash indexes:1-actor(act_id)
2-movie_cast(mov_id)
3-movie(mov_title)
Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

**********CONCLUSION QUERY 10 HASH INDEXES:**********

1-OPTIMIZED QUERY 10 HASH INDEXES COST = 16
2-OPTIMIZED QUERY 10 HASH INDEXES TIME = 0.6 MS

3)A)OLD QUERY 11 WITH NO INDEXES:

Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =off;

3-set enable_indexscan = off;
4- set enable_indexonlyscan = off;

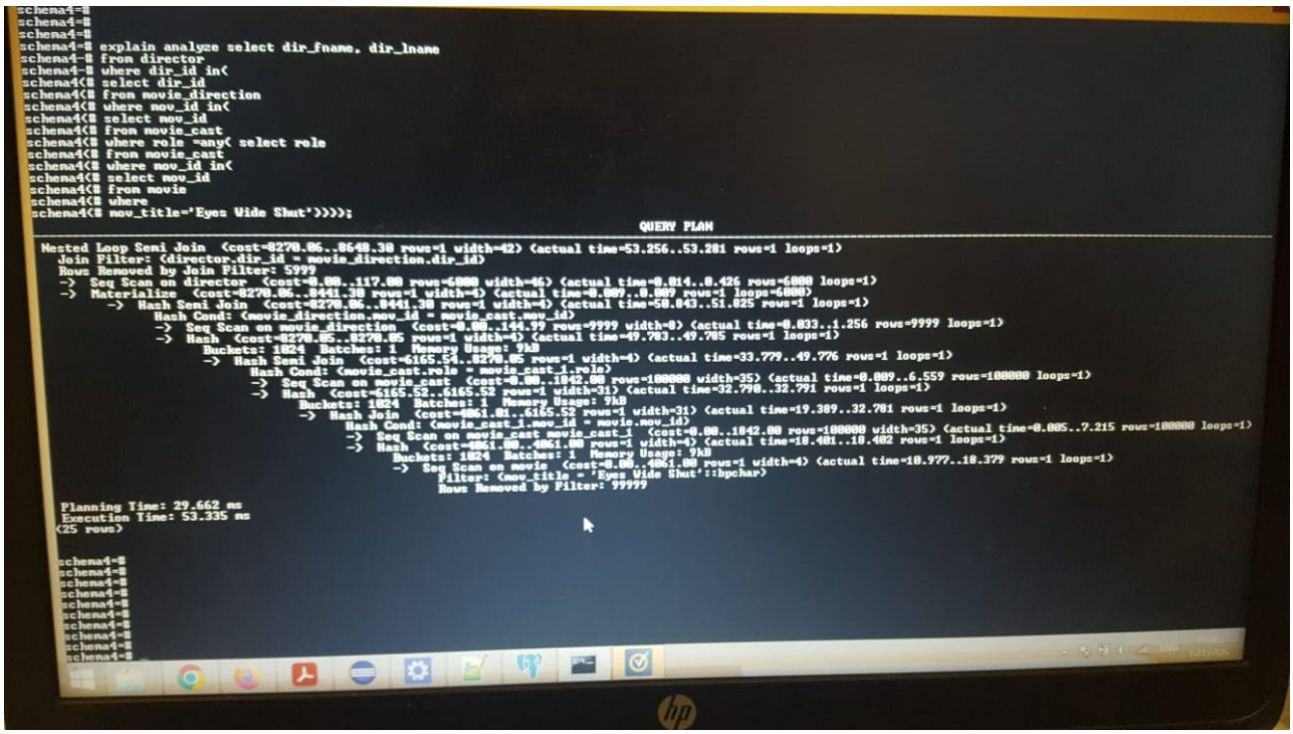

**********CONCLUSION OLD QUERY 11 NO INDEXES:**********

1-OLD QUERY 11 NO INDEXES COST = 8638
2-OLD QUERY 11 NO INDEXES TIME = 53 MS

3)B)OLD QUERY 11 WITH B+ INDEXES:

B+Indexes:
1- All default PKs indexes.
2 - movie(mov_title).
3 - movie_cast(mov_id).

4 - movie_cast(role).
5 -movie_direction(mov_id).
Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

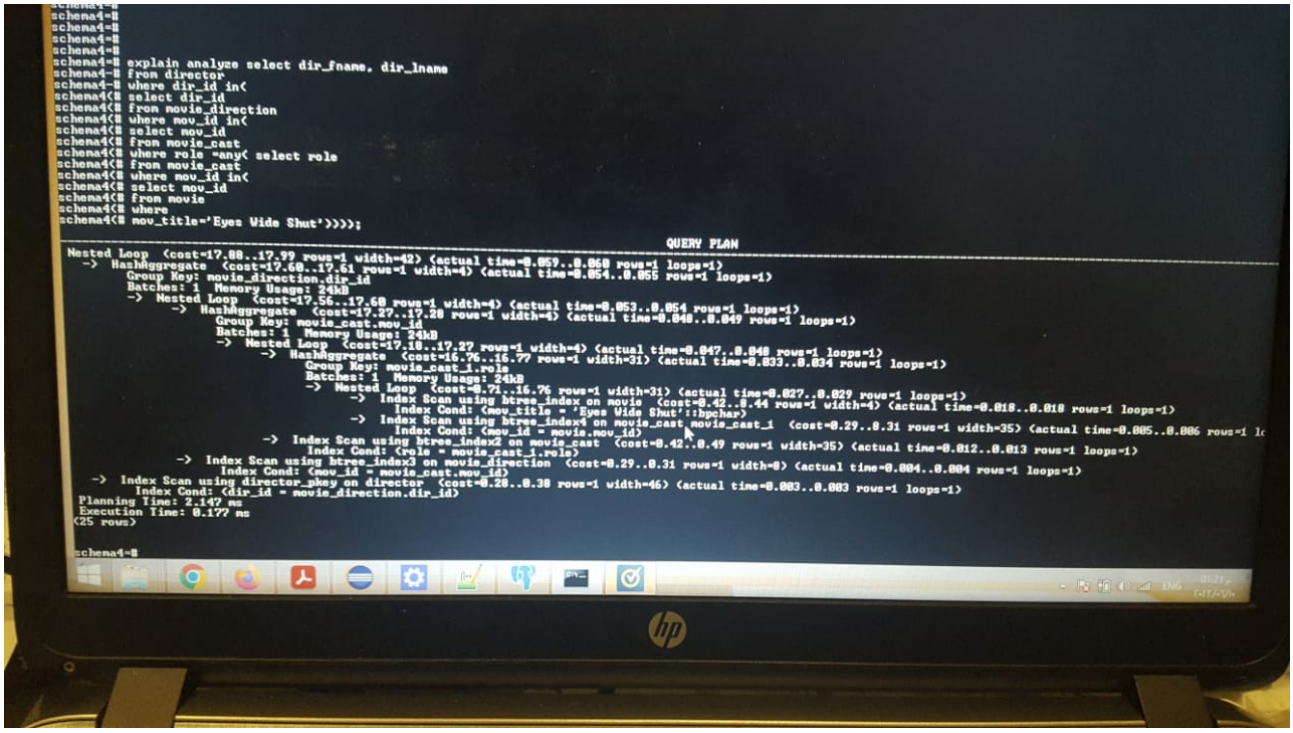

**********CONCLUSION OLD QUERY 11 B+ INDEXES:**********

1-OLD QUERY 11 NO INDEXES COST = 8638
2-OLD QUERY 11 NO INDEXES TIME = 53 MS

———————————————————————————

1-OLD QUERY 11 B+ INDEXES COST = 18
2-OLD QUERY 11 B+ INDEXES TIME = 0.1 MS
=> 99.8% more cost efficient & 99.8% more time efficient.

3)C)OLD QUERY 11 WITH Hash INDEXES:

Hash indexes: 1- director(dir_id)

2- movie_cast(mov_id)
3- movie_cast(role)
4 - movie_direction(mov_id)
5-movie(mov_title)

Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

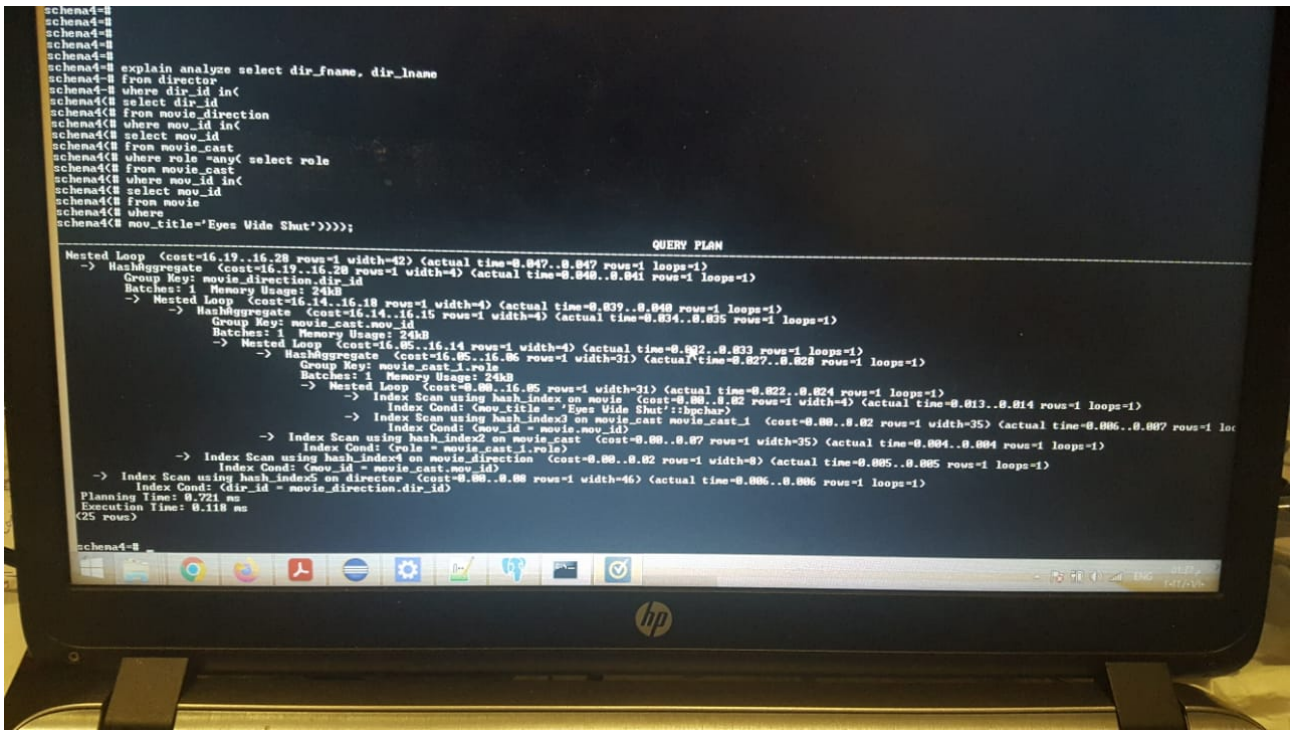

**********CONCLUSION QUERY 11 HASH INDEXES:**********

1-OLD QUERY 11 NO INDEXES COST = 8638
2-OLD QUERY 11 NO INDEXES TIME = 53 MS
3-OLD QUERY 11 HASH INDEXES COST = 16.2
4-OLD QUERY 11 HASH INDEXES TIME = 0.1 MS
=>99.8% more cost efficient &99.8% more time efficient.

3)D)OLD QUERY 11 WITH BRIN INDEX:

BRIN INDEXES: 1-movie(mov_title)
2- movie_cast(mov_id)

Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;

4- set enable_indexonlyscan = on;
5- update pg_index set indisvalid = false where indexrelid =
'director_pkey'::regclass ;
6- update pg_index set indisvalid = false where indexrelid =
‘movie_direction_pkey'::regclass ;

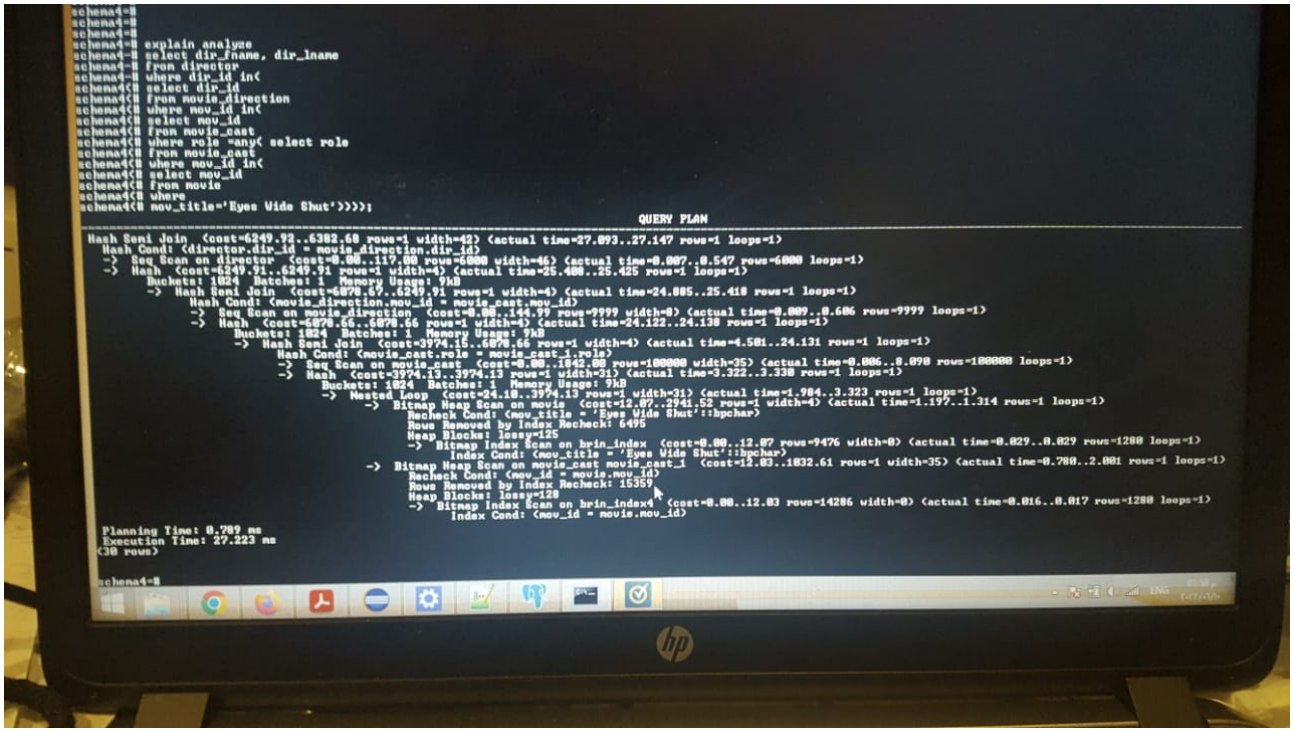

**********CONCLUSION OLD QUERY 11 BRIN INDEX:**********

1-OLD QUERY 11 NO INDEXES COST = 8638
2-OLD QUERY 11 NO INDEXES TIME = 53 MS.

———————————————————————

1-OLD QUERY 11 BRIN INDEX COST = 6382
2-OLD QUERY 11 BRIN INDEX TIME = 27 MS
=>26% more cost efficient &49% more time efficient.

3)E)OLD QUERY 11 WITH MY CHOICE FOR INDEXES:

Same as 3)c) using only hash indexes as all the conditions are equality conditions.

Hash indexes: 1- director(dir_id)
2- movie_cast(mov_id)
3- movie_cast(role)
4 - movie_direction(mov_id)
5-movie(mov_title)

Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

**********CONCLUSION OLD QUERY 11 HASH INDEXES:**********

1-OLD QUERY 11 NO INDEXES COST = 8638
2-OLD QUERY 11 NO INDEXES TIME = 53 MS
3-OLD QUERY 11 HASH INDEXES COST = 16.2
4-OLD QUERY 11 HASH INDEXES TIME = 0.1 MS
=>99.8% more cost efficient &99.8% more time efficient.

4)A)OPTIMIZED QUERY 11 WITH NO INDEXES:

Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =off;
3-set enable_indexscan = off;
4- set enable_indexonlyscan = off;

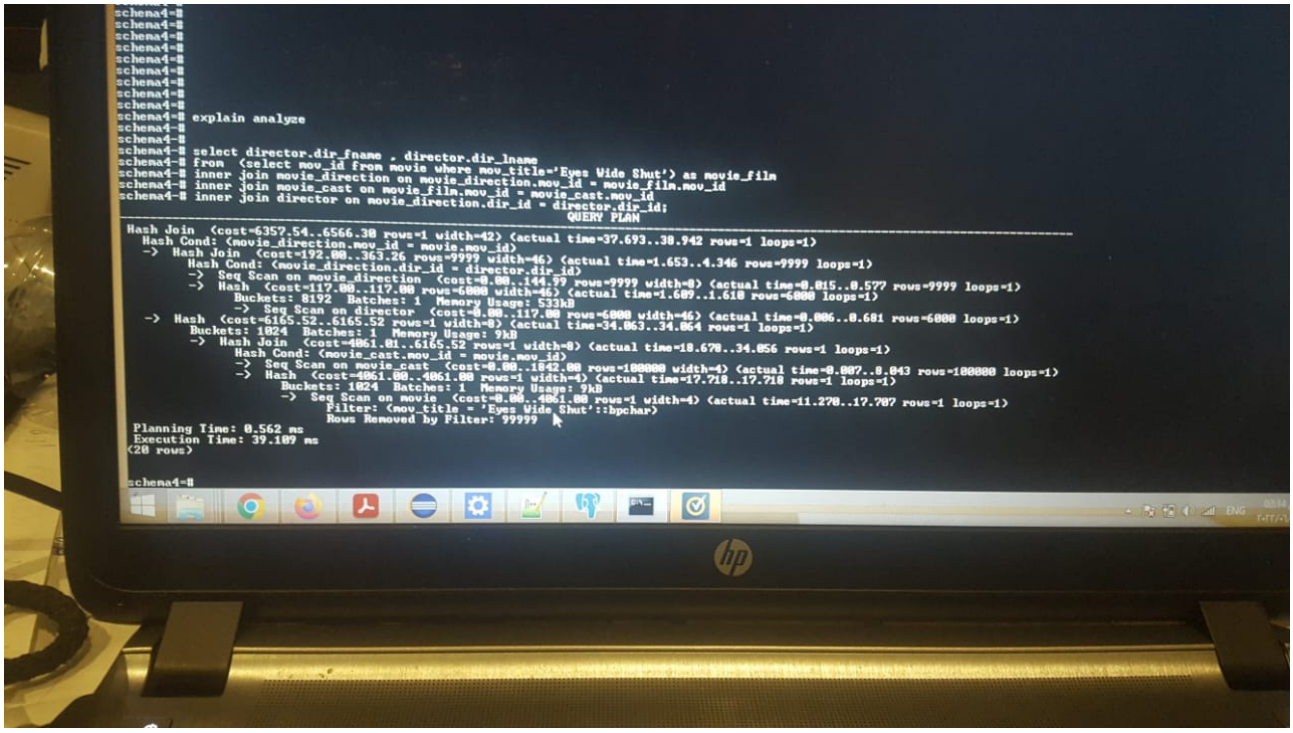

**********CONCLUSION OPTIMIZED QUERY 11 NO INDEXES:**********

1-OLD QUERY 11 NO INDEXES COST = 8638
2-OLD QUERY 11 NO INDEXES TIME = 53 MS
3-OPTIMIZED QUERY 11 NO INDEXES COST = 6566
4-OPTIMIZED QUERY 11 NO INDEXES TIME = 39 MS
=>24% more cost efficient &26% more time efficient.

4)B)OPTIMIZED QUERY 11 WITH B+ INDEXES:

B+Indexes:
1- All default PKs indexes.
2 - movie(mov_title).
3 - movie_cast(mov_id).
4 -movie_direction(mov_id).

Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

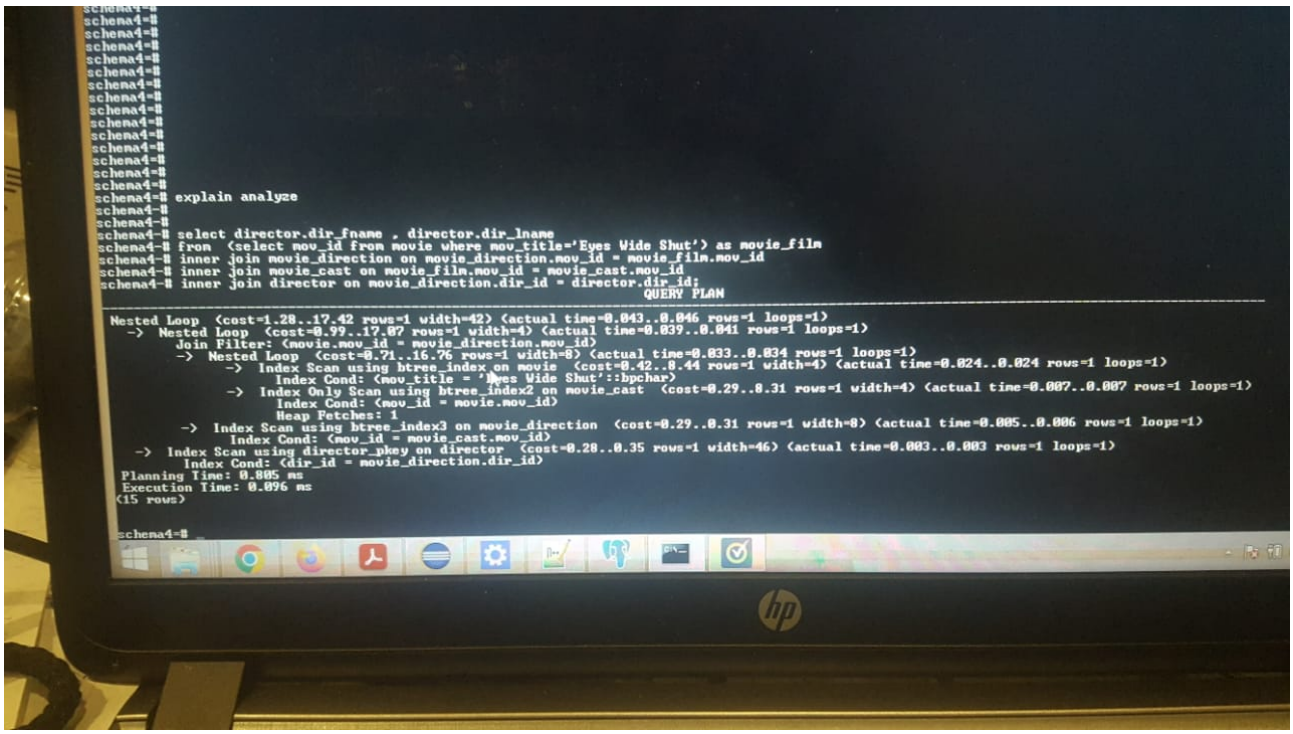

**********CONCLUSION OPTIMIZED QUERY 11 B+ INDEXES:**********

1-OLD QUERY 11 B+ INDEXES COST = 18
2-OLD QUERY 11 B+ INDEXES TIME = 0.1 MS
3-OPTMIZED QUERY 11 B+ INDEXES COST = 17
4-OPTIMIZED QUERY 11 B+ INDEXES TIME = 0.09 MS

=> 5.5% more cost efficient & 10% more time efficient.

4)C)OPTIMIZED QUERY 11 WITH Hash INDEXES:

Hash indexes: 1- director(dir_id)
2- movie_cast(mov_id)
3 - movie_direction(mov_id)
4-movie(mov_title)

Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

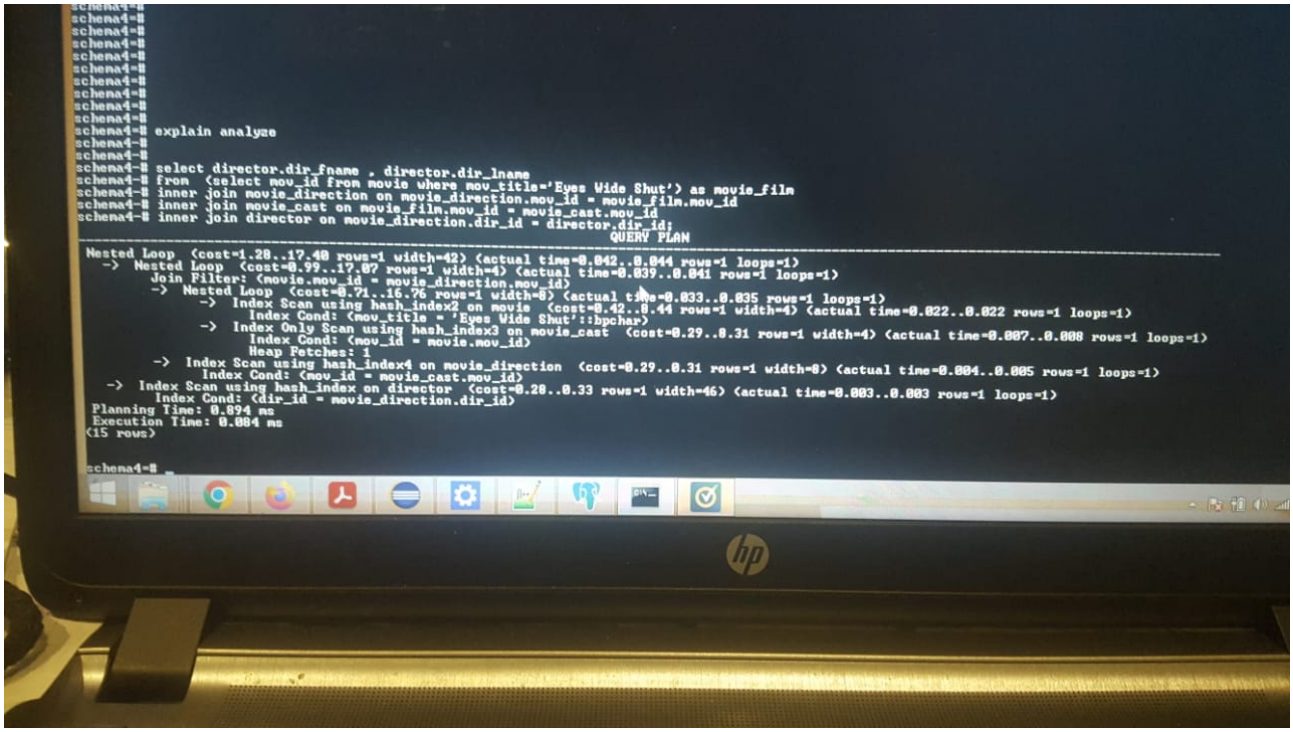

**********CONCLUSION OPTIMZED QUERY 11 HASH INDEXES:**********

1-OLD QUERY 11 HASH INDEXES COST = 16.2
2-OLD QUERY 11 HASH INDEXES TIME = 0.1 MS
3-OPTIMIZED QUERY 11 HASH INDEXES COST = 17
4-OPTMIZED QUERY 11 HASH INDEXES TIME = 0.084 MS

=> APPROX SAME COST & TIME

4)D)OPTIMIZED QUERY 11 WITH BRIN INDEX:

BRIN INDEXES: 1-movie(mov_title)
2- movie_cast(mov_id)
3- movie_direction(mov_id)

Flags: 1- set enable_seqscan = on
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;
5- update pg_index set indisvalid = false where indexrelid = ‘director_pkey'::regclass;
6- update pg_index set indisvalid = false where indexrelid = ‘movie_direction_pkey'::regclass ;
7- update pg_index set indisvalid = false where indexrelid = ‘movie_cast_pkey'::regclass ;

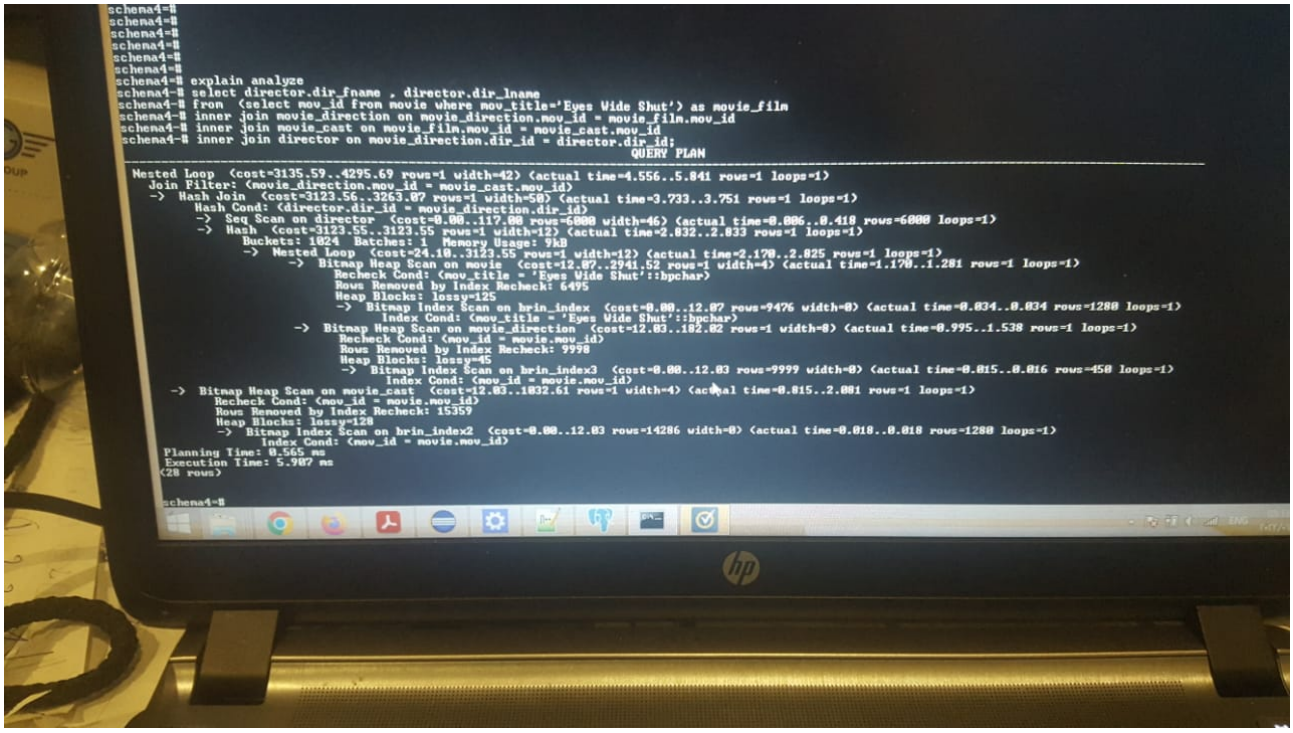

**********CONCLUSION OPTIMIZED QUERY 11 BRIN INDEX:**********
1-OLD QUERY 11 BRIN INDEX COST = 6382
2-OLD QUERY 11 BRIN INDEX TIME = 27 MS
3-OPTIMIZED QUERY 11 BRIN INDEX COST = 4295
4-OPTIMIZED QUERY 11 BRIN INDEX TIME = 6 MS
=>32.7% more cost efficient &77.7% more time efficient.
****

4)E)OPTIMIZED QUERY 11 WITH MY CHOICE FOR INDEXES:

Same as 4)c) using only hash indexes as all the conditions are equality conditions.

Hash indexes: 1- director(dir_id)
2- movie_cast(mov_id)
3 - movie_direction(mov_id)
4-movie(mov_title)
Flags: 1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

**********CONCLUSION OPTIMIZED QUERY 11 HASH INDEXES:**********
1-OLD QUERY 11 HASH INDEXES COST = 16.2
2-OLD QUERY 11 HASH INDEXES TIME = 0.1 MS
3-OPTIMIZED QUERY 11 HASH INDEXES COST = 17
4-OPTMIZED QUERY 11 HASH INDEXES TIME = 0.084 MS
=> APPROX SAME COST & TIME

5)A)OLD QUERY 12 WITH NO INDEXES:

Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =off;
3-set enable_indexscan = off;
4- set enable_indexonlyscan = off;

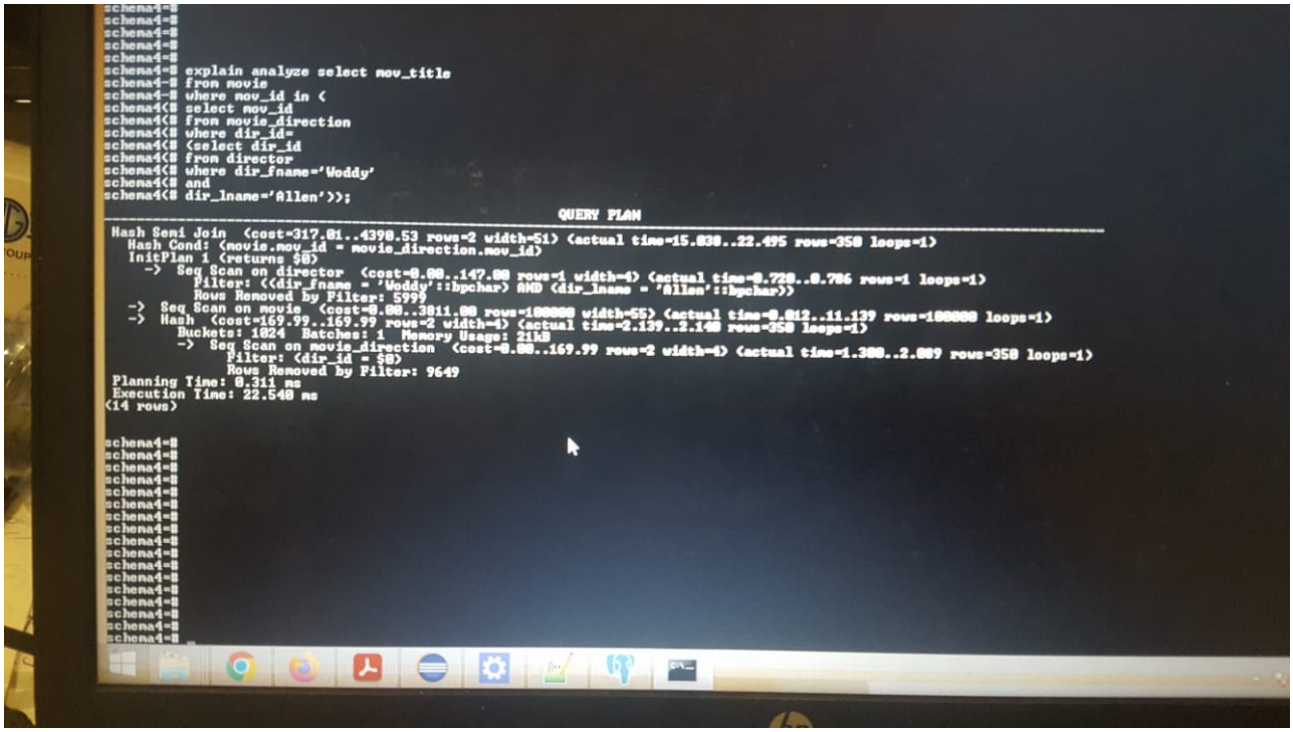

**********CONCLUSION OLD QUERY 12 NO INDEXES:**********

1-OLD QUERY 12 NO INDEXES COST = 4390
2-OLD QUERY 12 NO INDEXES TIME = 22.5 MS

5)B)OLD QUERY 12 WITH B+ INDEXES:

Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;

B+ Indexes: 1- movie_direction(dir_id)
2 - multi column B+ tree on director(dir_fname,dir_lname)
3- any default PKs B+ trees

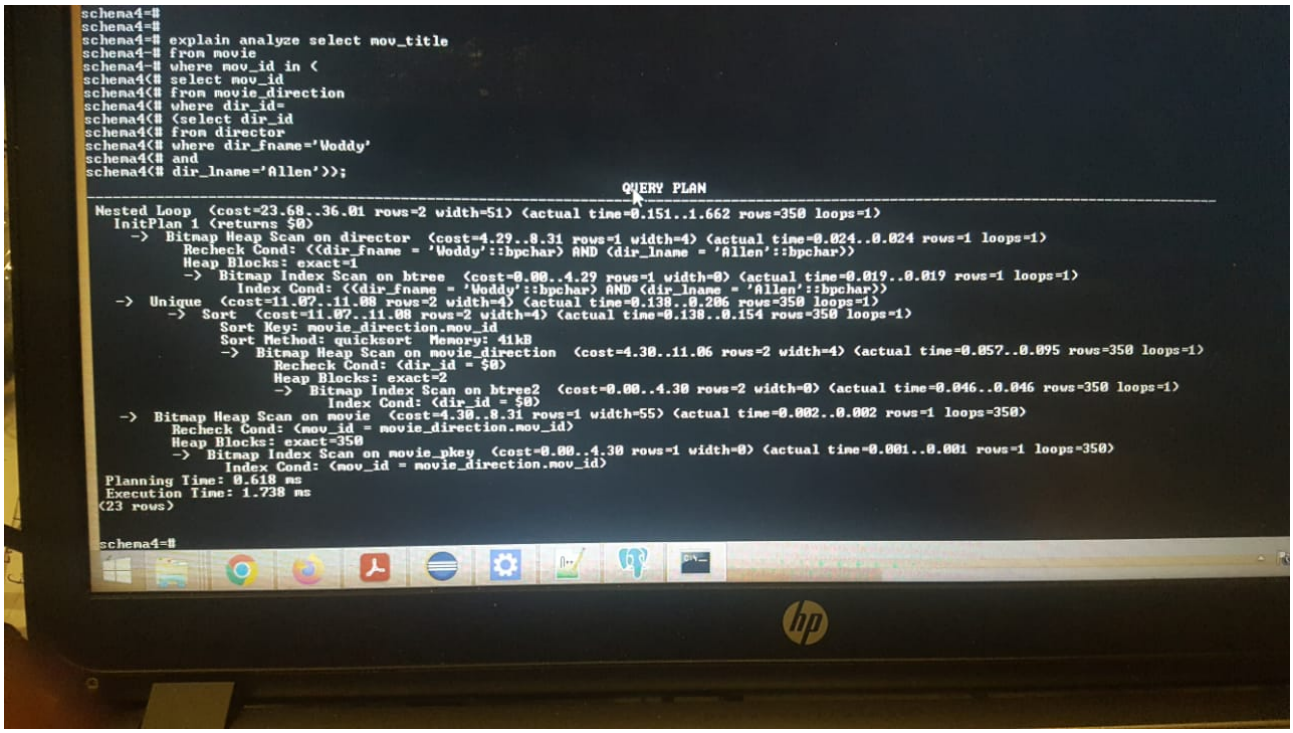

**********CONCLUSION OLD QUERY 12 B+ INDEXES:**********

1-OLD QUERY 12 NO INDEXES COST = 4390
2-OLD QUERY 12 NO INDEXES TIME = 22.5 MS

3-OLD QUERY 12 B+ INDEXES COST = 36
4-OLD QUERY 12 B+ INDEXES TIME = 1.7 MS

=>99% more cost efficient & 92.4 more time efficient

5)C)OLD QUERY 12 WITH HASH INDEXES:

Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;
Hash Indexes:1 - director(dir_lname)
2 - movie_direction(dir_id)
3 - movie(mov_id)
NOTE : Postgres doesn't support multi column indexes using Hash yet
So couldn't make a single hash on dir_lname & dir_fname
Also if a single hash table is created for dir_fname and another for dir_lname the query optimiser uses only one and fetches the other using filter but doesn't use both hash tables so I had to choose a single hash on dir_lname .

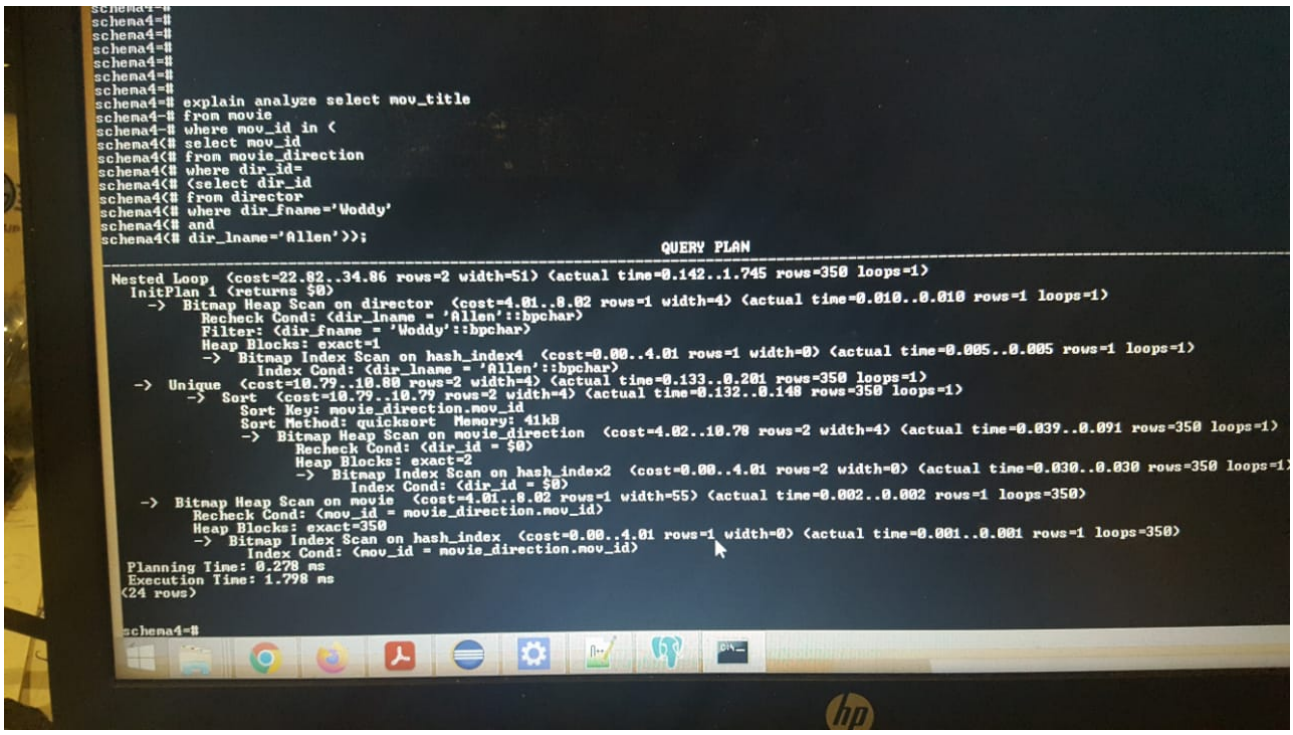

**********CONCLUSION OLD QUERY 12 Hash INDEXES:**********

1-OLD QUERY 12 NO INDEXES COST = 4390
2-OLD QUERY 12 NO INDEXES TIME = 22.5 MS
————————————————————————————————
3-OLD QUERY 12 Hash INDEXES COST = 35
4-OLD QUERY 12 Hash INDEXES TIME = 1.8 MS
=>99% more cost efficient & 92.4 more time efficient
****

5)D)OLD QUERY 12 WITH BRIN INDEX:

Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;
5- update pg_index set indisvalid = false where indexrelid = ‘movie_pkey’::regclass;
6- update pg_index set indisvalid = false where indexrelid = ‘movie_direction_pkey’::regclass;

BRIN Indexes: 1- movie(mov_id)
2 - movie_direction(dir_id)
3 - multi column BRIN director(dir_fname,dir_lname)

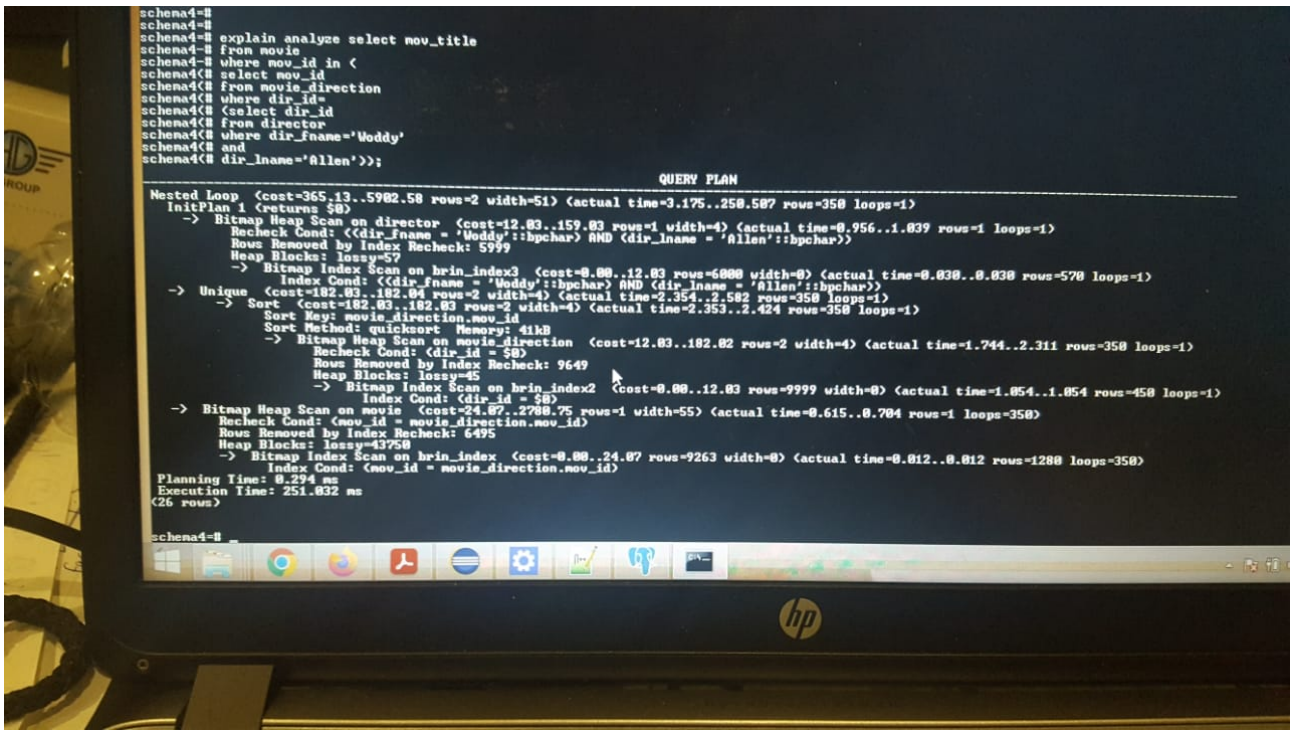

**********CONCLUSION OLD QUERY 12 BRIN INDEXES:**********

1-OLD QUERY 12 NO INDEXES COST = 4390
2-OLD QUERY 12 NO INDEXES TIME = 22.5 MS
————————————————————————————————
3-OLD QUERY 12 BRIN INDEXES COST = 5902
4-OLD QUERY 12 BRIN INDEXES TIME = 251 MS
=> Much less efficient cost and time wise compared to no indexes this is
This is mainly because BRINs should be on very large tables
(On the clustering columns) .

****

5)E)OPTIMIZED QUERY 12 WITH MY CHOICE FOR INDEXES:

I choose B+trees same as 5)b)
Flags:1- set enable_seqscan = on;
2- set enable_bitmapscan =on;
3-set enable_indexscan = on;
4- set enable_indexonlyscan = on;
B+ Indexes: 1- movie_direction(dir_id)
2 - multi column B+ tree on director(dir_fname,dir_lname)
3- any default PKs B+ trees
*OLD QUERY 12 B+ INDEXES COST = 36
*OLD QUERY 12 B+ INDEXES TIME = 1.7 MS

6)optimised query 12:

No better query than the old query as the join won’t give any additional benefits ,yes this is a nested query but it contains ZERO dependencies
 So it is a perfect nested query .